

# SPoP Project

Dr. Rahul Jain, rENIAC Inc

# Problem

- Build a In-Memory Key-Value Storage Software in C++
- Supported APIs
  - `get(key)`: returns value for the key
  - `put(key, value)`: add key-value, overwrite existing value
  - `delete(key)`
  - `get(int N)`: returns Nth key-value pair
  - `delete(int N)`: delete Nth key-value pair
- Spec
  - Max key size: 64 bytes
  - Each key char can be (a-z) or A-Z): Total 52 possible chars
  - Max Value Size: 256 bytes, any ASCII value
  - No DS/boost/STL etc Libraries to be used

# Class Definition

```
struct Slice{
    uint8_t size;
    char*    data;
};

class kvStore {
public:
    kvStore(uint64_t max_entries);
    bool get(Slice &key, Slice &value): //returns false if key didn't exist
    bool put(Slice &key, Slice &value): //returns true if value overwritten
    bool delete(Slice &key);
    bool get(int N, Slice &key, Slice &value): //returns Nth key-value pair
    bool delete(int N): //delete Nth key-value pair
};
```

# Evaluation Benchmark

- Multithreaded benchmark application
- Runs:
  - Benchmark would first load data via put calls (10 million entries, 2 min time limit)
  - Perform Single Threaded Transactions to verify kvStore functionality
  - **Multiple Transaction Threads with each thread calling one of the APIs**
- Evaluation Metrics:
  - TPS
  - CPU Usage
  - Memory Usage

# Relevant Concepts

- Ordered Data Structures
  - B-Tree, BST, etc
  - Tries, FST (Advanced)
  - Hybrid: Combination of Hash Table + Tree
  - .....
- Bit Hacks
- Cache Optimizations
- Memory Allocation Optimizations
  - Memory Object Reuse
  - No Dynamic Memory Allocation
- Multithreading

# Submission Deadline

- Form Groups of 2, can be different from assignment group
- Read papers or other relevant material around better data structures
  - must provide reference
- Implementation Spec: 3rd Feb
- Code+Report: 18th Feb **(Hard Deadline)**
- If a working code has been submitted on/before 18th Feb, a more optimized version can be submitted upto **29th Feb**